

VASYL'S PERSONAL PROJECTS

- **textstack.app — Open-Source AI Reader for Technical Books**

textstack.app · github.com/mrviduus/textstack

Local-first AI reading platform for finishing dense technical books. Local LLM via Ollama powers vocabulary generation, contextual explanations, and spaced-repetition card authoring. Stack: ASP.NET Core 10 + PostgreSQL 16 + React 19 + React Native (Expo). Validated under load at 63K requests / 100% success / p95 20.5 ms on a single container, with measured cost of \$0.002 per cycle (cache-aggressive strategy). AGPL-3.0. Live in production. Created independently: architecture, implementation, deployment, and writing.

- **AiAgents — 6 AI Agent Patterns in C# / .NET 10**

github.com/mrviduus/AiAgents

Learn-by-building solution covering the spectrum of agent patterns: BasicOpenAiExamples (fundamentals), ConsoleAgent (ReAct loop with weather tool), InvoiceAgentApi (single-agent RAG over documents), McpServer (MCP server exposing tools, prompts, and resources), RawJsonImplementation (low-level API). Multi-provider abstraction supporting OpenAI, Anthropic Claude, and Google Gemini. Docker-orchestrated.

- **rag-chatbot-dotnet — RAG with Pinecone + Semantic Kernel**

github.com/mrviduus/rag-chatbot-dotnet

Production-pattern RAG implementation. Indexes Wikipedia landmark articles into Pinecone via OpenAI embeddings, served through Microsoft Semantic Kernel chain. .NET 10. Documented architecture covering vector retrieval, top-K with system prompt and retrieved context injection.

- **ReplyMate — Local LLM Chrome Extension (WebLLM)**

github.com/mrviduus/ReplyMate

Chrome extension providing AI reply suggestions for LinkedIn, running entirely on-device via WebLLM. TypeScript. No server, no API key, no telemetry — on-device LLM inference in browser through WebGPU.

- **LoadSurge + xUnitOTel — Observability for AI Workloads**

github.com/mrviduus/LoadSurge · github.com/mrviduus/xUnitOTel

Open-source actor-based load testing framework (LoadSurge) and OpenTelemetry distributed tracing for xUnit v3 (xUnitOTel). Used in combination to validate the textstack inference pipeline under burst load.

- **Dev.to Technical Writing — @mrviduus**

dev.to/mrviduus

Production-focused technical posts. "I shipped local LLM features two months ago — production never ran them once" — Gemma 4 Challenge submission, post-mortem on shipping LLM features without monitoring. "I put Ollama on a 4 GB mobile GPU and got 2.5× — here's the VRAM math" — measurement of GPU offload trade-offs on commodity hardware.

CONTACTS

- **Email:** vasyl.vdov@gmail.com
- **LinkedIn:** linkedin.com/in/vasyl-vdovychenko
- **GitHub:** github.com/mrviduus
- **Personal blog:** vasyl.blog